

Fundamentals of Robot Vision

SGN-45006, 5 study credits

Esa Rahtu

Laboratory of Signal Processing, Tampere University

Model estimation (fitting)

- Least-squares
- Robust fitting
- RANSAC
- Hough transform

These topics are covered in Szeliski's book briefly, but more thoroughly in Chapter 17 of Forsyth & Ponce:


<http://courses.cs.washington.edu/courses/cse455/02wi/readings/book-7-revised-a-indx.pdf>

Acknowledgement: many slides from Svetlana Lazebnik, Derek Hoiem, Kristen Grauman, David Forsyth, Marc Pollefeys, and others (detailed credits on individual slides)

Relevant reading

- These topics are covered in Szeliski's book briefly, but more thoroughly in the following books:
 - Chapter 17 of Forsyth & Ponce:
<http://cmuems.com/excap/readings/forsyth-ponce-computer-vision-a-modern-approach.pdf>
 - Chapter 4 of Hartley & Zisserman:
[http://cvrs.whu.edu.cn/downloads/ebooks/Multiple%20View%20Geometry%20in%20Computer%20Vision%20\(Second%20Edition\).pdf](http://cvrs.whu.edu.cn/downloads/ebooks/Multiple%20View%20Geometry%20in%20Computer%20Vision%20(Second%20Edition).pdf)

Fitting




Fitting

- We've learned how to detect edges, corners, blobs. Now what?
- We would like to form a higher-level, more compact representation of the features in the image by grouping multiple features according to a simple model



Fitting

- Choose a *parametric model* to represent a set of features



simple model: lines



simple model: circles



complicated model: car

Fitting: Issues

Case study: Line detection



- **Noise** in the measured feature locations
- **Extraneous data:** clutter (outliers), multiple lines
- **Missing data:** occlusions

Fitting: Overview

- If we know which points belong to the line, how do we find the “optimal” line parameters?
 - Least squares
- What if there are outliers?
 - Robust fitting, RANSAC
- What if there are many lines?
 - Voting methods: RANSAC, Hough transform
- What if we're not even sure it's a line?
 - Model selection (not covered)


Least squares line fitting

Data: $(x_1, y_1), \dots, (x_n, y_n)$

Line equation: $y_i = m x_i + b$

Find (m, b) to minimize

$$E = \sum_{i=1}^n (y_i - m x_i - b)^2$$




Least squares line fitting

Data: $(x_1, y_1), \dots, (x_n, y_n)$

Line equation: $y_i = mx_i + b$

Find (m, b) to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$



$$E = \|Y - XB\|^2 \quad \text{where} \quad Y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} \quad X = \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \quad B = \begin{bmatrix} m \\ b \end{bmatrix}$$

$$E = \|Y - XB\|^2 = (Y - XB)^T (Y - XB) = Y^T Y - 2(XB)^T Y + (XB)^T (XB)$$

$$\frac{dE}{dB} = 2X^T XB - 2X^T Y = 0$$


$$X^T XB = X^T Y \quad \text{Normal equations: least squares solution to } XB = Y$$

Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines

Total least squares

Distance between point (x_i, y_i) and line $ax+by=d$ ($a^2+b^2=1$): $|ax_i + by_i - d|$




Total least squares

Distance between point (x_i, y_i) and line $ax+by=d$ ($a^2+b^2=1$): $|ax_i + by_i - d|$

Find (a, b, d) to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$




Total least squares

Distance between point (x_i, y_i) and line $ax+by=d$ ($a^2+b^2=1$): $|ax_i + by_i - d|$

Find (a, b, d) to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$

$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0$$



$$d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$


$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$

$$\frac{dE}{dN} = 2(U^T U)N = 0$$

Solution to $(U^T U)N = 0$, subject to $\|N\|^2 = 1$: eigenvector of $U^T U$ associated with the smallest eigenvalue (least squares solution to *homogeneous linear system* $UN = 0$)

Least squares: Robustness to noise

Least squares fit with an outlier:



Problem: squared error heavily penalizes outliers


Robust estimators

- General approach: find model parameters θ that minimize

$$\sum_i \rho(r_i(x_i, \theta); \sigma)$$


$r_i(x_i, \theta)$ – residual of i th point w.r.t. model parameters θ

ρ – robust function with scale parameter σ




The robust function ρ behaves like squared distance for small values of the residual u but saturates for larger values of u

Choosing the scale: Just right




The effect of the outlier is minimized

Choosing the scale: Too small



The error value is almost the same for every point and the fit is very poor

Choosing the scale: Too large



Behaves much the same as least squares

Robust estimation: Details


- Robust fitting is a nonlinear optimization problem that must be solved iteratively
- Least squares solution can be used for initialization
- Scale of robust function should be chosen adaptively based on median residual

RANSAC


- Robust fitting can deal with a few outliers – what if we have very many?
- Random sample consensus (RANSAC):
Very general framework for model fitting in the presence of outliers
- Outline
 - Choose a small subset of points uniformly at random
 - Fit a model to that subset
 - Find all remaining points that are “close” to the model and reject the rest as outliers
 - Do this many times and choose the best model

M. A. Fischler, R. C. Bolles. [Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography](#). Comm. of the ACM, Vol 24, pp 381-395, 1981.

RANSAC for line fitting example




RANSAC for line fitting example




Least-squares fit

RANSAC for line fitting example




1. Randomly select minimal subset of points

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. **Select points consistent with model**

RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop


RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop


RANSAC for line fitting example

Uncontaminated sample



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

RANSAC for line fitting

Repeat N times:

- Draw s points uniformly at random
- Fit line to these s points
- Find *inliers* to this line among the remaining points (i.e., points whose distance from the line is less than t)
- If there are d or more inliers, accept the line and refit using all inliers

Choosing the parameters

- Initial number of points **s**
 - Typically minimum number needed to fit the model
- Distance threshold **t**
 - Choose **t** so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2 = 3.84\sigma^2$
- Number of samples **N**
 - Choose **N** so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)

Choosing the parameters

- Initial number of points **s**
 - Typically minimum number needed to fit the model
- Distance threshold **t**
 - Choose **t** so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2 = 3.84\sigma^2$
- Number of samples **N**
 - Choose **N** so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$


s	proportion of outliers e						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

Choosing the parameters

- Initial number of points **s**
 - Typically minimum number needed to fit the model
- Distance threshold **t**
 - Choose **t** so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2 = 3.84\sigma^2$
- Number of samples **N**
 - Choose **N** so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$



Choosing the parameters


- Initial number of points **s**
 - Typically minimum number needed to fit the model
- Distance threshold **t**
 - Choose **t** so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2 = 3.84\sigma^2$
- Number of samples **N**
 - Choose **N** so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)
- Consensus set size **d**
 - Should match expected inlier ratio

Adaptively determining the number of samples

- Outlier ratio e is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield $e=0.2$
- Adaptive procedure:
 - $N=\infty$, $sample_count=0$
 - While $N > sample_count$
 - Choose a sample and count the number of inliers
 - If inlier ratio is highest of any found so far, set $e = 1 - (\text{number of inliers})/(\text{total number of points})$
 - Recompute N from e :
$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$
 - Increment the $sample_count$ by 1

RANSAC pros and cons

- Pros
 - Simple and general
 - Applicable to many different problems
 - Often works well in practice
- Cons
 - Lots of parameters to tune
 - Doesn't work well for low inlier ratios (too many iterations, or can fail completely)
 - Can't always get a good initialization of the model based on the minimum number of samples




Fitting: Review

- Least squares
- Robust fitting
- RANSAC

Fitting: Review

- ✓ If we know which points belong to the line, how do we find the “optimal” line parameters?
 - ✓ Least squares
- ✓ What if there are outliers?
 - ✓ Robust fitting, RANSAC
- What if there are many lines?
 - Voting methods: RANSAC, Hough transform

Fitting: The Hough transform



Voting schemes

- Let each feature vote for all the models that are compatible with it
- Hopefully the noise features will not vote consistently for any single model
- Missing data doesn't matter as long as there are enough features remaining to agree on a good model

Hough transform

- An early type of voting scheme
- General outline:
 - Discretize *parameter space* into bins
 - For each feature point in the image, put a vote in every bin in the parameter space that could have generated this point
 - Find bins that have the most votes






Image space




Hough parameter space

P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959


Parameter space representation

- A line in the image corresponds to a point in Hough space




Parameter space representation

- What does a point (x_0, y_0) in the image space map to in the Hough space?




Parameter space representation

- What does a point (x_0, y_0) in the image space map to in the Hough space?
 - Answer: the solutions of $b = -x_0m + y_0$
 - This is a line in Hough space





Parameter space representation

- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?



Parameter space representation

- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?
 - It is the intersection of the lines $b = -x_0m + y_0$ and $b = -x_1m + y_1$




Parameter space representation

- Problems with the (m, b) space:
 - Unbounded parameter domains
 - Vertical lines require infinite m

Parameter space representation

- Problems with the (m,b) space:
 - Unbounded parameter domains
 - Vertical lines require infinite m
- Alternative: *polar representation*



Each point (x,y) will add a sinusoid in the (θ, ρ) parameter space

Algorithm outline

- Initialize accumulator H to all zeros
- For each feature point (x, y) in the image

For $\theta = 0$ to 180

$$\rho = x \cos \theta + y \sin \theta$$


$$H(\theta, \rho) = H(\theta, \rho) + 1$$

end


end

- Find the value(s) of (θ, ρ) where $H(\theta, \rho)$ is a local maximum
- The detected line in the image is given by
$$\rho = x \cos \theta + y \sin \theta$$


H : accumulator array (votes)



Basic illustration



features




votes


[Hough transform demo](#)

Other shapes



Square




Circle




Several lines



A more complicated image




Effect of noise




features

Effect of noise



features




votes


Peak gets fuzzy and hard to locate

Effect of noise


- Number of votes for a line of 20 points with increasing noise:



Random points



features




votes

Uniform noise can lead to spurious peaks in the array

Random points

- As the level of uniform noise increases, the maximum number of votes increases too:




Dealing with noise

- Choose a good grid / discretization
 - **Too coarse:** large votes obtained when too many different lines correspond to a single bucket
 - **Too fine:** miss lines because some points that are not exactly collinear cast votes for different buckets
- Increment neighboring bins (smoothing in accumulator array)
- Try to get rid of irrelevant features
 - E.g., take only edge points with significant gradient magnitude

Incorporating image gradients

- Recall: when we detect an edge point, we also know its gradient direction
- But this means that the line is uniquely determined!
- Modified Hough transform:



$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

For each edge point (x,y)

θ = gradient orientation at (x,y)

$\rho = x \cos \theta + y \sin \theta$

$H(\theta, \rho) = H(\theta, \rho) + 1$


end

Hough transform for circles


- How many dimensions will the parameter space have?
- Given an unoriented edge point, what are all possible bins that it can vote for?
- What about an *oriented* edge point?

Hough transform for circles

image space




Hough parameter space



Hough transform for circles

- Conceptually equivalent procedure: for each (x, y, r) , draw the corresponding circle in the image and compute its “support”



Is this more or less efficient than voting with features?

Review: Hough transform

- Hough transform for lines
- Hough transform for circles
- Hough transform pros and cons

Hough transform: Pros and cons

- Pros
 - Can deal with non-locality and occlusion
 - Can detect multiple instances of a model
 - Some robustness to noise: noise points unlikely to contribute consistently to any single bin
- Cons
 - Complexity of search time increases exponentially with the number of model parameters
 - Non-target shapes can produce spurious peaks in parameter space
 - It's hard to pick a good grid size