

What we'd like to do?

- Visual scene understanding
- What is in the image and where
- Object categories, identities, properties, activities, relations, ...

Things vs Stuff

- Thing: object with specific shape
- Stuff: material defined by homogeneous or repetitive pattern. Has no specific spatial extent or shape
  - ↳ Segmentation method

Recognition tasks

- Image classification: does the image contain an aeroplane?
- Object class detection/localization: where are the aeroplanes if any?
- Object class segmentation: which pixels are part of an aeroplane?
- Panoptic segmentation: Besides object segmentation, also background segmentation

Challenges

- Background clutter: lot of stuff in the background
- Occlusions and truncation: partially seen objects from the camera
- Intra-class variation: a class can have ~~are~~ lot of different versions
  - ↳ object instant recognition: recognizes a specific model of an object  
not its generic class

category detection: recognizes generic classes  
it is harder to perform than instant detection

So why bother?

- Spatial relationships for image understanding and retrieval ("a cat riding a skateboard")
- Visual question and answering: object grasping/tracking  
(collision prevention, face recognition)

## Sliding window detectors

### Problem of background clutter : Solution

- Use sub window:
  - At correct position, no clutter is present
  - Slide window to detect objects
  - Change size of the window to search over scales

### Detection by classification

- Basic component : binary classifier
  - ↳ sliding window over window using CNNs is too slow
- Detect objects in clutter by searching
  - sliding window : exhaustive search over position and scale
    - ↳ in practice, it is possible to use same window size over spatial pyramid pooling
    - ↳ more efficient

### Window (image) classification

- Features usually engineered
- Classifier learned from data



### Problems with sliding windows


- Aspect ratio
- Granularity (finite grid)
- Partial occlusions
- Multiple responses



(Conference on CV and PR)


- Objective: detect (localize) standing humans in an image
- Sliding window classifier
- Train a binary SVM classifier on whether a window contains a standing person or not
- Histogram of Oriented Gradients (HOG) feature
- Although introduced for pedestrian detection, it has been successfully used with many other categories

Window (image) classifier



Feature: HOG

- Tile  $64 \times 128$  pixel window into  $8 \times 8$  pixel cells
- Calculate gradient image (edge detection)
- Each cell represented by histogram over 8 orientation bins (or sector)



- Adds a second level of overlapping spatial bins re-normalizing orientation histogram over larger spatial area

- Feature dimensions (approx) =  $16 \times 8$  (for tiling)  $\times 8$  (orientations)  $\times 4$  (blocks) = 4096

- Similarity to CNN

{ CNN learns the filters automatically  
 { Sum pooling  
 { normalization

- OK job

## Linear classifier

-  $f(x) = w^T x + b$

2D discriminant is a line

- It learns such weights  $w$  and bias  $b$  that

3D " is a plane

$$f(x_i) = \begin{cases} \geq 0 & y_i = 1 \\ < 0 & y_i = -1 \end{cases}$$

## Linear separability


- The points might be linearly separable but with very narrow margin
- The large margin solution might be better, even constraints are violated
- In general there is a trade off between the margin and the number of mistakes on the training data

## Support Vector Machine

- It is a way to optimize this trade off
- Find a good trade off between the classification margin and the misclassified training points

$$f(x) = w^T x + b$$

$$\min_{w \in \mathbb{R}^d} \|w\|^2 + C \sum_i \max(0, 1 - y_i f(x_i))$$



Learned Model using HOG detector  
→ evidence it is a person

- Positive weights
  - Negative weights
- Average over positive training data
- evidence it is not a person

$$f(x) = w^T x + b$$

→ comes from 8x8

What do negative weights mean?

OS. 3

$$\left. \begin{aligned} w^T x &> 0 \\ (w_+ - w_-) x &> 0 \\ w_+ x &> w_- x \end{aligned} \right\}$$

- Complete system compete pedestrian / pillar / doorway models
- Discriminative models come with own background model
- Avoid detections on doorways by penalizing vertical edges

Pedestrian model > Pedestrian background model

Problems when training a sliding window detector

- Inherently asymmetric problem: many more "non-object" than "objects"
- Classifier needs to have very low false positive rate
- "Non-object" category is very complex and needs a lot of data

Optimizing approach: Boosting

1. Pick a negative training set at random
2. Train ~~the~~ classifier
3. Run on training data
4. Add false positives to training set
5. Repeat from step 2. (Retraining)

Data augmentation: With you available data you

- Flip
  - Rotate
  - Scale
  - Crop
  - Translate
  - Apply Gaussian noise
- Jittered positive  $\uparrow$   
to the images  
 $\downarrow$   
Jittered

! Window (image) first stage classification



Jittered positive and random negative

Feature extraction  
HOG

$$x \rightarrow \begin{bmatrix} \vdots \\ \vdots \\ \vdots \end{bmatrix} \rightarrow \text{Linear SVM classifier} \\ f(x) = w^T x + b$$

Precision - Recall curve


- Precision: % of returned windows that are correct
- Recall: % of correct windows that are returned





## Evaluating the detected bounding boxes

- Area of overlap (AO) measure
- Correct detection if intersection over union larger than threshold



$$AO(B_{gt}, B_p) = \frac{|B_{gt} \cap B_p|}{|B_{gt} \cup B_p|}$$

Second training phase → Retrain using better data


- Find high scoring false positive detections
- Use them as hard negatives for next training round
- Cost = # training image · inference time per image

## First training phase

Accelerating sliding window search

- Sliding window search is slow since some many windows are needed
- $M \times N \times \text{scale} = 100\,000$  windows for  $320 \times 240$  image
- Most windows are negative
- It is possible to speed up the search


First training phase



Second training phase



## Cascade classification



In the end, you spend much less computational cost

- Slow and expensive classifiers only applied to a few windows
- Controlling complexity vs speed: numbers of features, numbers of parts, ...

Detection proposal: Hierarchically clustering superpixels

- larger homogeneous area is considered a unit in the image  $\rightarrow$  superpixel

- Hierarchical segmentation: start with small and merge using mcs

Produces roughly 2000 regions per image with 95% of hitting relevant objects

Things to remember

- Detection by sliding window classification  $\rightarrow$  Concept and components

- Multiple scales (and aspect ratios) to detect objects of different size

- Importance of hard negative mining (due to class imbalance)

- Cascade detectors  $\rightarrow$  Speed up inference

- Speed up training and inference by selecting sub-set of windows only