

Image features

- Low-level features (corners, edges, textures patches) are needed for many tasks:
 - Image matching and registration
 - Structure-from-motion and image-based 3D modelling
 - Image segmentation
 - Object recognition
 - Image retrieval

~~Mid-level features~~

Edge detection

- Goal: Identify sudden changes (discontinuities) in an image
 - Intuitively, most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
- Ideal: artist's line drawing

Why do we care about edges?

- Extract information, recognize objects
 - Recover geometry and viewpoint
- } robust against illumination changes




Origin of edges

Caused by:

- surface normal discontinuity
- depth discontinuity
- surface color discontinuity
- illumination discontinuity

Edge detection

- An edge is a place of rapid change in the image intensity function



edges correspond to extreme of derivative

Derivative with convolution

For 2D $f(x, y)$: $\frac{\partial f(x, y)}{\partial x} = \lim_{\epsilon \rightarrow 0} \frac{f(x+\epsilon, y) - f(x, y)}{\epsilon}$

For discrete data, infinite differences: $\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$

To implement this as convolution, this is the associated filter:

$$\begin{array}{ccc} 0 & 0 & 0 \\ 0 & \boxed{-1 \quad 1} & 0 \\ 0 & 0 & 0 \end{array}$$

$$\frac{\partial f(x, y)}{\partial x}$$

$$\begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

$$\frac{\partial f(x, y)}{\partial y}$$

Finite difference filters

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$

$$M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$

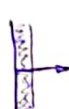
$$M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$


Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$

$$M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

Image gradient

The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

 $\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$

 $\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

$$\theta = \arctan \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right) : \text{Gradient direction}$$


The gradient points in the direction of most rapid increase in intensity

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2} : \text{Gradient strength}$$

Application: Gradient - domain image editing

- Goal: solve for pixel values in the target region to match gradients of the source region while keeping background pixels the same

Effects of noise




Derivative theorem of convolution

- Differentiation is convolution, and convolution is associative:

$$\frac{\partial}{\partial x} (f * g) = f * \frac{\partial}{\partial x} g$$

- This saves us an operation:



$f * \frac{\partial}{\partial x} g$



Solution: smooth first

$f(x)$: signal



g : Kernel



$f * g$: convolution



$\frac{\partial}{\partial x} (f * g)$: differentiation



Separability of the Gaussian filter

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) = \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right) \right)$$

So if smoothing you apply

Smoothed derivative removes noise, but blurs edges. Also finds edges at different scales

Smoothing vs derivative filters



- Smoothing filters

- Gaussian: remove high frequency components: low-pass filter
- Their values can be negative, but usually all values are positive
- All values should sum 1, so that constant regions are not affected



- Derivative filters

- Derivatives of Gaussian
- They can be negative
- The values should sum 0
- No response in constant regions





Building edge detector

1° Filters in x and y direction and then take the responses into the gradient vector \rightarrow magnitude and orientation \Leftarrow norm of the gradient

2° Thresholded norm of the gradient \rightarrow to discard weak edges

How to turn thick curves into actual lines?

3° Non-maximum suppression



Check if pixel is local maximum along gradient direction, select single max across width of the edge
- requires checking interpolated pixels p and n

Some edges then disappear to be too weak despite being relevant. How to solve, lowering threshold? No

4° Hysteresis thresholding

You do not apply the same threshold everywhere

Use a high threshold to start edge curves and low threshold to continue them

All of this is wrapped up in Canny edge detector

MATLAB [edge(image, 'canny')]

Data-driven edge detection

Edge detection can work as ground truth for feeding a ^{ML classifier} ~~classifier~~ instead of labelling manually every edge in every image.
The output should like the edge detection output but not like a binary output. Real value between 0 and 1.

Keypoint extraction: Corners

Why extract keypoints?

- Panorama stitching: We have 2 images, how do we combine them?
- o Step 1: extract points
- o Step 2: match keypoint features
- o Step 3: align images

Characteristics of good keypoints


- Repetability: same keypoint can be found in several images despite geometric and photometric transformations
- Saliency: each keypoint is distinctive
- Compactness and efficiency: many fewer keypoints than image pixels
- Locality: a keypoint occupies a relatively small area of the image; robust to clutter and occlusion

Applications of keypoints


- Image alignment
- 3D reconstruction
- Motion tracking
- Robot navigation
- Indexing and database retrieval
- Object recognition

Corner detection: Basic idea


- We should easily recognize the point by looking through a small window
- Shifting a window in any direction should give a large change in intensity



"flat" region:
no change in
all directions



"edge":
no change along
the edge detection



"corner":
significant change
in all directions

Corner detection: mathematics

- Change in appearance of window W for the shift $[u, v]$:

$$E(u, v) = \sum_{(x, y) \in W} [I(x+u, y+v) - I(x, y)]^2$$

- First-order Taylor approximation for small motions $[u, v]$:

$$I(x+u, y+v) \approx I(x, y) + I_x u + I_y v$$

$$E(u, v) = \sum_{(x, y) \in W} [I(x+u, y+v) - I(x, y)]^2 \approx$$

$$\approx \sum_{(x, y) \in W} [I(x, y) + I_x u + I_y v - I(x, y)]^2 = \sum_{(x, y) \in W} [I_x u + I_y v]^2 =$$

$$= \sum_{(x, y) \in W} [I_x^2 u^2 + I_y^2 v^2 + 2 I_x I_y u v]$$

$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$ where M is a second moment matrix obtained from image derivatives

$$M = \begin{bmatrix} \sum_{x, y} I_x^2 & \sum_{x, y} I_x I_y \\ \sum_{x, y} I_x I_y & \sum_{x, y} I_y^2 \end{bmatrix}$$

the sums are over all the pixels in the window

$E(u, v)$




circle if perfectly symmetric
otherwise, ellipse

↳ horizontal size $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

Diagonalization of M

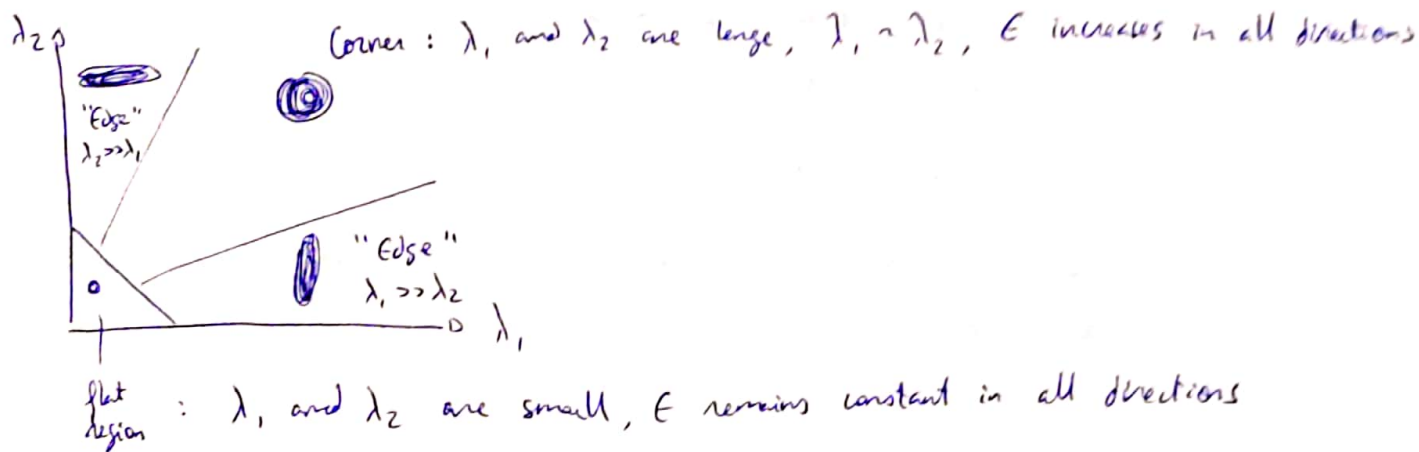
$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$



Axis lengths determined by eigenvalues
Orientation determined by R

Interpreting the eigenvalues

4



Corner response function

$$R = \det(M) - \alpha \text{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2 \begin{cases} > 0 : \text{corner} \\ < 0 : \text{NOT a corner} \\ = 0 : \text{flat region} \end{cases}$$

Harris Corner detector

1. Compute partial derivatives at each pixel
2. Compute second moment matrix M in a Gaussian window around each pixel
3. Compute corner response function R
4. Threshold R
5. Find local maxima of response function (non-maximum suppression)

Robustness of corner features

- The scale of intensities matter
- Partially invariant to affine intensity scale
- Image translation: Corner location is covariant with translation
- " rotation: " " " " rotation
- Scaling: it is not invariant



SIFT Keypoint detection


- CNN are much better for semantic problems
- Good enough for many applications

Keypoint detection with scale selection

- We want to extract keypoints with characteristic scale that is invariant with the image transformation

Basic idea

- Blob detection filter
- Convolve the image with a "blob filter" at multiple scales and look for extreme of filter response in the resulting scale space







- Find maxima and minima of blob filter response in space and scale

Blob filter

Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D


$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$



edge in zero crossing of second derivative


From edges to blobs

- Edge = ripple
- Blob = superposition of two ripples
- Spatial selection: the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is matched to the scale of the blob




Scale selection

- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases




Scale normalization

- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases



- To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
- Laplacian is the second derivative of the Gaussian so it must be multiplied by σ^2



Blob detection 2D

- Scale-normalized Laplacian of Gaussian

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

- At what scale does this achieve the maximum response to a binary circle of radius r ?

- The zeros of the Laplacian have to be aligned with the circle

- The Laplacian is given by (up to scale):

$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2 + y^2)/2\sigma^2}$$

$$\xrightarrow{\text{maximum}} \sigma = \frac{1}{\sqrt{2}} r$$



Scale-space blob detector

- 1 - Convolve image with scale-normalized Laplacian at several scales
- 2 - Find maxima of squared Laplacian response in scale-space

Eliminating edge responses

- Laplacian has strong response along edge
- Solution: filter based on Harris response function over neighborhoods containing the blobs

Efficient implementation

- Approximating the Laplacian with a difference of Gaussians:

$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma)) \quad \leftarrow \text{Laplacian}$$

$$\text{DOG} = G(x, y, k\sigma) - G(x, y, \sigma) \quad \leftarrow \text{Difference of Gaussians}$$

From feature detection to feature description

- Scaled and rotated versions of the same neighborhood will give rise to blobs that are related by the same transformation
- What to do if we want to compare the appearance of these image regions?
 - Normalization: transform these regions into same-size circles
 - Problem: rotational ambiguity

Eliminating rotational ambiguity

- To assign a unique orientation to circular image windows
 - Create histogram of local gradient directions in the patch
 - Assign canonical orientation at peak of smoothed histogram

SIFT descriptor computation

6

- Use the normalized region about the keypoint
- Compute gradient magnitude and orientation at each point in the region
- weight them by a Gaussian window overlaid on the circle
- create an orientation histogram over 4×4 subregions of the window
- 4×4 descriptors over 16×16 sample array were used in practice
 4×4 times 8 directions gives a vector of 128 values

Summary of SIFT feature detection and description

1 - Scale - space extreme detection

Search over multiple scales and image rotations

2 - Keypoint localization

Fit a model to determine location and scale

Select keypoints based on a measure of stability

3 - Orientation assignment

Compute best orientation(s) for each keypoint region

4 - Keypoint description

Use local image gradients at selected scale and rotation to describe each keypoint region

It is robust to image condition changes

Properties of SIFT

Extraordinarily robust detection and description technique

- Can handle changes in viewpoint
 - up to about 60 degree out-of-plane rotation
- Can handle significant changes in illumination (day vs night sometimes)
- Fast and efficient - can run in real time
- Lots of code available