

Visual classification (CIFAR-10 dataset)

Be prepared for the exercise sessions. You may ask TA questions regarding your solutions, but don't expect them to show you how to start from the scratch. Before the end of the session, demonstrate your solution to TA to receive exercise points.

1. **CIFAR-10 – Bayesian classifier** (40 points)

We continue with the CIFAR-10 dataset and we will adopt Bayesian classifier approach this time.

First we need to select suitable features. Let's use the **average color of each image as their imagery feature**. Instead of the $x = 32 \times 32 \times 3 = 3072$ length pixel vector **for each image, you need** to convert them to $\mathbf{f} = (m_R, m_G, m_B)$ values where m_R is the mean value of the red channel in the image, m_G green and m_B blue, i.e. **every image is represented with three values**.

Let's assume that channels are independent and normal distributed, i.e. Bayesian probability can be computed from

$$P(class_1|\mathbf{f}) = \frac{P(\mathbf{f}|class_1)P(class_1)}{P(\mathbf{f}|class_1)P(class_1) + P(\mathbf{f}|class_2)P(class_2) + \dots}.$$

Note that **you may omit the demoninator parts since its same for all classes**. Now,

$$P(\mathbf{f}|class_1)P(class_1) = \mathcal{N}(m_R; \mu_{R,c_1}, \sigma_{R,c_1})\mathcal{N}(m_G; \mu_{G,c_1}, \sigma_{G,c_1})\mathcal{N}(m_B; \mu_{B,c_1}, \sigma_{B,c_1})P(c_1)$$

where μ and σ are the mean and variance of each class and for each feature.

Write a function `f=cifar_10_features(x)` that forms the mean color feature of \mathbf{x} . Write another function `[mu,sigma,p]=cifar_10_bayes_learn(F,labels)` that computes the normal distribution parameters for the samples in \mathbf{F} (one line per sample) and respective labels in `labels`. Note that **the size of mu and sigma must be `numberofclasses` \times 3**. The `p` (prior) probabilities are a vector of **`numberofclasses` \times 1**, one for each class.

Finally write function `c=cifar_10_bayes_classify(f,mu,sigma,p)` that returns the Bayesian optimal class `c` for the sample `f`.

Run your classifier for all CIFAR-10 test samples and report the accuracy (write another script for this).

Hints: `mean()`, `std()`, `normpdf()`

2. **CIFAR-10 – Bayesian classifier with better pdf** (20 points)

In this experiment we continue the previous Bayesian classifier, but we relax the naive assumption and **switch to the full *multivariate normal distribution*** (see the `mvnpdf` function in Matlab). That means that **you need to put all three features** (image mean of the Red, Green and Blue channels) **into a single feature** vector $\mathbf{f}_{1 \times 3}$ and represent class specific pdf's using the mean vectors $\boldsymbol{\mu}_{3 \times 1}$ and covariance $\boldsymbol{\Sigma}_{3 \times 3}$.

Compute the classification accuracy and compare it to the naive version - which one is better and why?

3. **CIFAR-10 – Bayesian with extended features** (40 points)

Perhaps you can improve your classification with better features?

In the previous exercises you used three mean values for the whole 32×32 images as a feature $\mathbf{f}_{3 \times 1}$. However, you may divide each image to four 16×16 sub-images which provide $4 \times$ more features $\rightarrow \mathbf{f}_{12 \times 1}$. You may continue the process and divide each sub-image to another four 8×8 sub-sub-image providing now $16 \times$ more features $\rightarrow \mathbf{f}_{48 \times 1}$. Note that the image division to sub-images can be done in multiple ways. You can see the sub-images as 4 squares forming the original image, or 4 rows or columns side by side forming the image. The way the image i.e. the pixels are divided to the sub-images affects the results slightly.

Extend your feature extraction function `f=cifar_10_features(x,N)` such that now N defines the sub-window size (32 would be the original) and the function returns these new advanced *spatial features*. Test at least values $N = 32, 16, 8, 4$ and plot a graph where accuracy is plotted as a function of the sub-window size. Use full multivariate normal distribution in your experiments.

What kind of behaviour you will find and what explains that?

Hints: `mvnpdf()`, `cov()`