

SGN-41007 Pattern Recognition and Machine Learning

Exercise Set 6: February 11–February 15, 2019

Exercises consist of both pen&paper and computer assignments. Pen&paper questions are solved at home before exercises, while computer assignments are solved during exercise hours. The computer assignments are marked by `python` and Pen&paper questions by `pen&paper`

1. `pen&paper` Count the number of parameters in a neural network

Consider the traditional shallow neural network architecture of Figure 1. Suppose our inputs are 64×64 RGB bitmaps of two categories of traffic signs.

Let the network structure be the following:

- The input is $64 \times 64 \times 3 = 12288$ -dimensional
- On the 1st layer there are 100 nodes (marked in blue)
- On the 2nd layer there are 100 nodes (marked in blue)
- On the 3rd (output) layer there are 10 nodes (marked in blue; one for each class)

Compute the number of parameters (coefficients) in the net.




Figure 1: Vanilla neural network.

2. **pen&paper** Consider the following Keras code defining a convolutional neural network.

```
N = 10          # Number of feature maps
w, h = 5, 5     # Conv. window size

model.add(Conv2D(N, (w, h),
                  input_shape=(64, 64, 3),
                  activation = 'relu',
                  padding = 'same'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(N, (w, h),
                  activation = 'relu',
                  padding = 'same'))
model.add(MaxPooling2D((2, 2)))

model.add(Flatten())
model.add(Dense(2, activation = 'sigmoid'))
```

- a) Draw a diagram of the network similar to the one at the bottom of slide 14 in <http://www.cs.tut.fi/courses/SGN-41007/slides/Lecture6.pdf>
- b) Compute the number of parameters of the network at each layer (and explain why).
3. **python** Load Traffic sign data for deep neural network processing.

Download an extended version of the two class German Traffic Sign Recognition Benchmark (GTSRB) dataset from

http://www.cs.tut.fi/courses/SGN-41007/GTSRB_subset_2.zip

This time, images are in color and there are about 400 from both classes.

After collecting the data, normalize all samples into range [0,1]; i.e., subtract `numpy.min(X)` and divide the result by `numpy.max(X)`.

Finally, split the data to training and testing (80% / 20%) using `sklearn.cross_validation.train_test_split`.

4. `python` Define the network in Keras.

Edit the network of Question 2 in your code such that `model.summary()` gives the following output:

<code>model.summary()</code>		
<hr/>		
<code>Layer (type)</code>	<code>Output Shape</code>	<code>Param #</code>
<hr/>		
<code>conv2d_49 (Conv2D)</code>	<code>(None, 64, 64, 32)</code>	<code>2432</code>
<hr/>		
<code>max_pooling2d_47 (MaxPooling)</code>	<code>(None, 16, 16, 32)</code>	<code>0</code>
<hr/>		
<code>conv2d_50 (Conv2D)</code>	<code>(None, 16, 16, 32)</code>	<code>25632</code>
<hr/>		
<code>max_pooling2d_48 (MaxPooling)</code>	<code>(None, 4, 4, 32)</code>	<code>0</code>
<hr/>		
<code>flatten_15 (Flatten)</code>	<code>(None, 512)</code>	<code>0</code>
<hr/>		
<code>dense_29 (Dense)</code>	<code>(None, 100)</code>	<code>51300</code>
<hr/>		
<code>dense_30 (Dense)</code>	<code>(None, 2)</code>	<code>202</code>
<hr/>		
<code>Total params: 79,566</code>		
<code>Trainable params: 79,566</code>		
<code>Non-trainable params: 0</code>		
<hr/>		

5. `python` *Compile and train the net.*

Compile and train the network following the examples of the lecture slides and documentation at <http://keras.io/>.

Use the following parameters:

- **Loss:** categorical_crossentropy (same thing as log loss; see previous exercises)
- **Optimizer:** stochastic gradient descent
- **Minibatch size:** 32
- **Number of epochs:** 20

Also add the parameter `metrics=['accuracy']` as an argument of `model.compile` and give the test data to training algorithm `model.fit(..., validation_data = [X_test, y_test])`. Then, the optimizer will report the test error every epoch.